

AMS 205B Take-Home Final Exam
Mickey Warner

Problem 1

(a)

The likelihood is given by

$$\begin{aligned} L(\mathbf{y}|\mathbf{x}, \alpha, \beta, \omega, \delta, \sigma^2) &= \prod_{i=1}^n (2\pi\sigma^2)^{-1/2} \exp \left\{ -\frac{1}{2\sigma^2} [y_i - \alpha - \beta \cos(\omega x_i + \delta)]^2 \right\} \\ &= (2\pi\sigma^2)^{-n/2} \exp \left\{ -\frac{1}{2\sigma^2} \sum_{i=1}^n [y_i - \alpha - \beta \cos(\omega x_i + \delta)]^2 \right\} \end{aligned}$$

If I take the partial derivative w.r.t. ω of the log-likelihood, we obtain

$$\frac{\partial}{\partial \omega} \log L = -\frac{\beta}{\sigma^2} \sum_{i=1}^n (y_i - \alpha - \beta \cos(\omega x_i + \delta)) \sin(\omega x_i + \delta) x_i$$

Thus solving for ω when this equals 0 has no available closed form. This further means that the m.l.e. for $\theta = (\alpha, \beta, \omega, \delta, \sigma^2)$ is not available in closed form.

To obtain estimates for the parameters, I use profile likelihoods and iteratively optimize (ω, δ) and $(\alpha, \beta, \sigma^2)$ using the `optim` function in R. The optimization is improved with some reasonable constraints (given our model and data):

$$\alpha \in (\min_i y_i, \max_i y_i)$$

$$\beta \in (\min_i y_i, \max_i y_i)$$

$$\omega \in (0, 3)$$

$$\delta \in [-\pi, \pi)$$

$$\sigma^2 \in (0.001, \text{var}(\mathbf{y}))$$

The estimates are $\hat{\theta} = (\hat{\alpha}, \hat{\beta}, \hat{\omega}, \hat{\delta}, \hat{\sigma}^2) = (1.969, 1.499, 1.752, -0.037, 0.019)$. The 3 in the bound for ω is a rough estimate for the maximum number of periods observed in the data. The code is given in the appendix.

(b)

To obtain an approximate confidence interval from a Wald-like test, we need to ensure that the regularity conditions are met.

- A1: our model assumes we have a random sample
- A2: given the constraint $\delta \in [-\pi, \pi)$, the model is identifiable
- A3: the parameters do not determine the support
- A4: there may be a concern if $\delta = -\pi$, but we could just as easily defined $\delta \in [0, \pi)$ to avoid this, so the parameter space contains an open set around the true parameters
- A5: the normal density is infinitely differentiable as is the function $\cos(\omega x_i + \delta)$ and the derivatives are continuous.
- A6: though I don't calculate the third derivative of the log, based on the second derivative, shown next,

The Wald approximations for ω and δ are given by

$$\frac{\hat{\omega} - \omega}{1/\sqrt{I(\hat{\omega})}} \sim N(0, 1), \quad \frac{\hat{\delta} - \delta}{1/\sqrt{I(\hat{\delta})}} \sim N(0, 1)$$

where $\hat{\omega}$ and $\hat{\delta}$ are the mles from above,

$$\begin{aligned} I(\hat{\omega}) &= -\frac{\partial^2}{\partial \omega^2} \log L(\mathbf{y}|\mathbf{x}, \alpha, \beta, \omega, \delta, \sigma^2) \Big|_{\theta=\hat{\theta}} \\ &= \frac{\beta}{\sigma^2} \sum_{i=1}^n ([y_i - \alpha - \beta \cos(\omega x_i + \delta)] \cos(\omega x_i + \delta) x_i^2 + [\beta x_i^2 \sin(\omega x_i + \delta)^2]) \Big|_{\theta=\hat{\theta}} \\ &= 281704.5 \end{aligned}$$

and,

$$\begin{aligned} I(\hat{\delta}) &= -\frac{\partial^2}{\partial \delta^2} \log L(\mathbf{y}|\mathbf{x}, \alpha, \beta, \omega, \delta, \sigma^2) \Big|_{\theta=\hat{\theta}} \\ &= \frac{\beta}{\sigma^2} \sum_{i=1}^n ([y_i - \alpha - \beta \cos(\omega x_i + \delta)] \cos(\omega x_i + \delta) + [\beta \sin(\omega x_i + \delta)^2]) \Big|_{\theta=\hat{\theta}} \\ &= 5397.7. \end{aligned}$$

These lead to the following 95% confidence intervals

$$\begin{aligned} \{\omega : \hat{\omega} - 1/\sqrt{I(\hat{\omega})} z_{0.975} < \omega < \hat{\omega} + 1/\sqrt{I(\hat{\omega})} z_{0.975}\} &= (1.748, 1.756) \\ \{\delta : \hat{\delta} - 1/\sqrt{I(\hat{\delta})} z_{0.975} < \delta < \hat{\delta} + 1/\sqrt{I(\hat{\delta})} z_{0.975}\} &= (-0.063, -0.010) \end{aligned}$$

(c)

The parametric bootstrap procedure yields the confidence intervals

$$\begin{aligned}\omega &\in (1.745, 1.760) \\ \delta &\in (-0.092, 0.016)\end{aligned}$$

which are wider than those based on the normal approximation.

(d)

I would test $H_0 : \delta = 0$ versus $H_a : \delta \neq 0$ based on the confidence interval from the bootstrap sample since this will be closer to the exact interval than the approximation would be. Since 0 is contained the interval, we do not have enough evidence to reject the null that $\delta = 0$.

Problem 2

(a)

The likelihood is

$$L \equiv L(\mathbf{x}, \mathbf{y} | \lambda_1, \lambda_2) = \lambda_1^{-n} \lambda_2^{-n} e^{-\sum x_i / \lambda_1} e^{-\sum y_i / \lambda_2}$$

The unconstrained maximum is found by taking the derivative of the log-likelihood and setting it equal to 0.

$$\frac{\partial}{\partial \lambda_1} \log L = -\frac{n}{\lambda_1} + \frac{1}{\lambda_1^2} \sum x_i \stackrel{set}{=} 0$$
$$\hat{\lambda}_1 = \bar{x}$$

The second derivative evaluated at $\lambda_1 = \hat{\lambda}_1$ gives

$$\left. \frac{\partial^2}{\partial \lambda_1^2} \log L \right|_{\lambda=\bar{x}} = \frac{n}{(\bar{x})^2} - \frac{2}{(\bar{x})^3} \sum x_i = -\frac{n^3}{(\sum x_i)^2} < 0$$

so we have a maximum. By similar logic, $\hat{\lambda}_2 = \bar{y}$.

Under the null hypothesis, the maximum depends on \bar{x} and \bar{y} . If $\bar{x} \leq \bar{y}$, then the constrained maximum occurs at the same location as the unconstrained. For the case $\bar{x} > \bar{y}$, we must check along the border $\lambda_1 = \lambda_2$. The other boundaries are not of interest since L will be increasing toward $\lambda_1 = \lambda_2$.

The mle under the constrained likelihood solves the equation

$$\begin{aligned}\frac{\partial}{\partial \lambda} \log L(\mathbf{x}, \mathbf{y} | \lambda_1 = \lambda, \lambda_2 = \lambda) &= -\frac{2n}{\lambda} + \frac{1}{\lambda^2} \left(\sum x_i + \sum y_i \right) \stackrel{set}{=} 0 \\ \Rightarrow \hat{\lambda} &= \frac{\sum x_i + \sum y_i}{2n}\end{aligned}$$

The likelihood ratio is then given by

$$\begin{aligned}\kappa \equiv LRT &= \begin{cases} 1 & \bar{x} \leq \bar{y} \\ \frac{\sup_{\lambda_1 \leq \lambda_2} L(\mathbf{x}, \mathbf{y} | \lambda_1, \lambda_2)}{\sup_{\lambda_1, \lambda_2} L(\mathbf{x}, \mathbf{y} | \lambda_1, \lambda_2)} & \bar{x} > \bar{y} \end{cases} \\ &= \begin{cases} 1 & \bar{x} \leq \bar{y} \\ \frac{(\bar{x}\bar{y})^n}{[\frac{1}{2}(\bar{x} + \bar{y})]^{2n}} & \bar{x} > \bar{y} \end{cases}\end{aligned}$$

We reject null when $\kappa < c$ for some constant $0 < c < 1$. We will reject only if $\bar{x} > \bar{y}$, so this is the case we're interested in. We simplify the rejection region

$$\begin{aligned}R &= \left\{ (x, y) : \frac{(\bar{x}\bar{y})^n}{[\frac{1}{2}(\bar{x} + \bar{y})]^{2n}} < c \right\} = \left\{ (x, y) : \frac{\bar{x}\bar{y}}{(\bar{x} + \bar{y})^2} < c \right\} \\ &= \left\{ (x, y) : \frac{(\bar{x} + \bar{y})^2}{\bar{x}\bar{y}} > c \right\} \\ &= \left\{ (x, y) : \frac{\bar{x}^2 + 2\bar{x}\bar{y} + \bar{y}^2}{\bar{x}\bar{y}} > c \right\} \\ &= \left\{ (x, y) : \frac{\bar{x}}{\bar{y}} + \frac{\bar{y}}{\bar{x}} + 2 > c \right\} \\ &= \left\{ (x, y) : \frac{\sum x_i}{\sum y_i} + \frac{\sum y_i}{\sum x_i} + 2 > c_\alpha \right\}\end{aligned}$$

where c_α is chosen so $Pr((\mathbf{x}, \mathbf{y}) \in R | \lambda_1 \leq \lambda_2) = \alpha$. The exact p -value is

$$Pr \left(\frac{\sum x_i}{\sum y_i} + \frac{\sum y_i}{\sum x_i} + 2 > c_\alpha; \lambda_1 \leq \lambda_2 \right)$$

It was at this point that I started thinking I was doing everything wrong. I have no idea how to compute this correctly. The closest I got is some gamma-looking distribution. I tried showing that we could reject if $\sum x_i + \sum y_i$ is high, but couldn't figure that out.

(b)

Since we're working with the exponential distribution, the regularity conditions are met and a Wald test for the statistic $T = \bar{X} - \bar{Y}$ can be given by

$$Z = \frac{\bar{X} - \bar{Y}}{1/\sqrt{I(\hat{\lambda})}} \sim N(0, 1)$$

where $I(\hat{\lambda})$ is the observed information number under the null (or rather, for $\lambda_1 = \lambda_2$). Again, something seemed wrong, especially given the p -values I calculated.

(c)

Under the null, the observations from both samples can be re-arranged in any order. We combine the x_i 's and y_i 's into a single vector, then randomly permute the elements. This first n are treated as the x_i 's and the last n the y_i 's. We then compute the (permuted) statistic $\bar{x} - \bar{y}$ of the permuted sample. We do this many times and store all of the statistics from the permuted sample. We can then compute a p -value by counting how many of the permuted statistics were larger than our observed and divide by the number of permutations made.

(d)

I wasn't able to get a definitive p -value for (a). My attempts gave varied results.

For (b), I calculated $p_B = 0.000389$, which seems outrageously small.

For (c), I got $p_C = 0.037$ which seems reasonable. Assuming I did anything right (which is doubtful), the p -values do not agree. I would learn toward p_C being the most correct, though it may be off since I didn't do all possible permutations. My conclusion would be that $\lambda_1 > \lambda_2$, but what do I know.

Problem 3

(a)

The power function is

$$\begin{aligned}\beta_1(\theta) &= P\left(\sqrt{n}\frac{(\bar{X} - \theta_0)}{S} > c_1\right) \\ &= P\left(\sqrt{n}\frac{(\bar{X} - \theta)}{S} > c_1 + \sqrt{n}\frac{(\theta_0 - \theta)}{S}\right) \\ &= P\left(T_{n-1} > c_1 - \sqrt{n}\frac{(\theta - \theta_0)}{S}\right)\end{aligned}$$

where T_{n-1} is a t random variable with $n - 1$ degrees of freedom and c_1 is chosen such that $\beta_1(\theta_0) = P(T_{n-1} > c_1) = \alpha$.

(b)

Assuming $T = \sqrt{n}(\bar{X} - \theta_0)/S$ follows a standard normal under the null hypothesis, an approximate power function is

$$\beta_2(\theta) = P\left(Z > c_2 - \sqrt{n}\frac{(\theta - \theta_0)}{S}\right)$$

where c_2 satisfies $\beta_2(\theta_0) = P(Z > c_2) = \alpha$. For T to be normal under the null, this would seem to imply T is also normal for $\theta > \theta_0$ (given our sample population).

(c)

Figure 1 shows graphs for $\beta_1(\theta)$ and $\beta_2(\theta)$ at $n = 10, n = 100$, and $n = 1000$. As n increases, the power increases, notably for θ close to θ_0 . This makes sense because if the true θ is different than θ_0 , however small, we'd expect to have more power in our test as the sample size increases. Also, β_2 is greater than β_1 for all θ .

(d)

I think the LRT could shown to reject for small values and since the hypotheses could be written as union-intersection tests, Theorem 8.3.21 could apply leaving us $\beta_2(\theta)$ as most powerful.

(e)

I take “true standard deviation of the data” to mean S . I think it should be σ , but based on my power functions, the correct things don't cancel. We need to find n that satisfies

$$P(Z > c_2) = \alpha \text{ and } \beta_2(\theta_0 + 2S) = P\left(Z > c_1 - \sqrt{n}\frac{(\theta_0 + 2S - \theta_0)}{S}\right) = 0.8$$

For $\alpha = 0.05$, $c_2 = 1.644$. This leaves us with $0.8 = P(Z > 1.644 - 2\sqrt{n})$. Since $P(Z > -0.8416) = 0.8$, $1.644 - 2\sqrt{n} = -0.8416$ implies $n = 1.545$, but we choose it to be the next largest integer so $n = 2$.

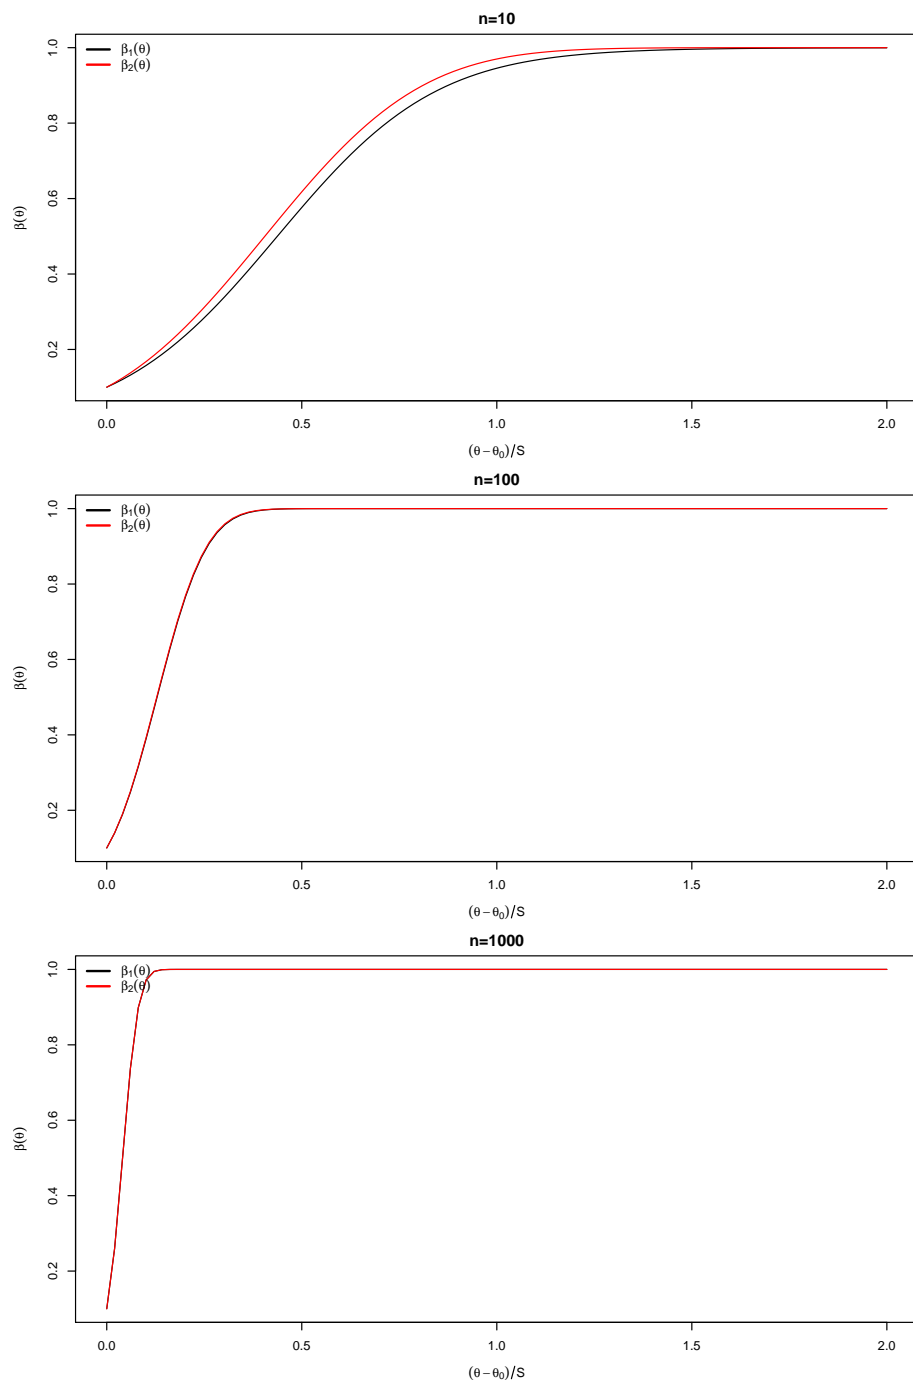


Figure 1: Graphs for the two power functions given three different sample sizes.

Some of the R code

```
orbits = read.table("~/files/data/205b/orbits.txt", header = TRUE)
n = NROW(orbits)

f1 = function(p, log = TRUE){
  alpha = p[1]
  beta = p[2]
  sig2 = p[3]
  out = -n/2 * log(2*pi*sig2) - 1/(2*sig2)*sum((orbits$Y - alpha -
    beta*cos(omega.hat * orbits$X + delta.hat))^2)
  if (log)
    return (out)
  return (exp(out))
}

f2 = function(p, log = TRUE){
  omega = p[1]
  delta = p[2]
  out = -n/2 * log(2*pi*sig2.hat) - 1/(2*sig2.hat)*sum((orbits$Y - alpha.hat -
    beta.hat*cos(omega * orbits$X + delta))^2)
  if (log)
    return (out)
  return (exp(out))
}

f3 = function(p, log = TRUE){
  alpha = p[1]
  beta = p[2]
  omega = p[3]
  delta = p[4]
  sig2 = p[5]
  out = -n/2 * log(2*pi*sig2) - 1/(2*sig2)*sum((orbits$Y - alpha -
    beta*cos(omega * orbits$X + delta))^2)
  if (log)
    return (out)
  return (exp(out))
}

### (a)
### Profile likelihood
# Initial values for the "fixed" parameters
alpha.hat = mean(c(range(orbits$Y))) # Estimate for the intercept
beta.hat = diff(range(orbits$Y))/2 # Estimate for amplitude
sig2.hat = 0.1 # Guess for variance

# The starting values for the optimizer
xy.1 = expand.grid(
  mean(c(range(orbits$Y))) + seq(-0.2, 0.2, length = 10),
  diff(range(orbits$Y))/2 + seq(-0.2, 0.2, length = 10),
  seq(0.005, 0.1, length = 10))
xy.2 = expand.grid(seq(0, 3, length = 40), seq(-pi, pi, length = 40))

# Iterate through the optimization 5 times (though only 2 may be really necessary)
for (j in 1:5){
  temp.2 = Inf
  for (i in 1:nrow(xy.2)){
    # Optimize treating alpha, beta, sig^2 as fixed (i.e. using alpha.hat,
    # beta.hat, and sig2.hat as the fixed values)
    temp = optim(as.double(xy.2[i,]), function(x) -f2(x, log = TRUE),
      method = "L-BFGS-B", lower = c(0, -pi), upper = c(3, pi))
    # If the i'th starting points produced a better mode, update the parameters
    if (temp$value < temp.2){
      temp.2 = temp$value
      omega.hat = temp$par[1]
      delta.hat = temp$par[2]
    }
  }
  temp.2 = Inf
  for (i in 1:nrow(xy.1)){
    # Optimize treating omega, delta as fixed (using omega.hat and delta.hat)
    temp = optim(as.double(xy.1[i,]), function(x) -f1(x, log = TRUE),
      method = "L-BFGS-B", lower = c(min(orbits$Y), min(orbits$Y), 0.001),
      upper = c(max(orbits$Y), max(orbits$Y), var(orbits$Y)))
    # If the i'th starting points produced a better mode, update the parameters
    if (temp$value < temp.2){
      temp.2 = temp$value
      alpha.hat = temp$par[1]
      beta.hat = temp$par[2]
      sig2.hat = temp$par[3]
    }
  }
}

### (b)
### Wald-like
l.omega = beta.hat / sig2.hat * sum(
  ((orbits$Y - alpha.hat - beta.hat*cos(omega.hat*orbits$X + delta.hat)) *
```



```

cos(omega.hat*orbits$X + delta.hat)*orbits$X^2) +
(beta.hat * orbits$X^2 * sin(omega.hat * orbits$X + delta.hat)^2))

I.delta = beta.hat /sig2.hat * sum(
((orbits$Y - alpha.hat - beta.hat*cos(omega.hat*orbits$X + delta.hat)) *
cos(omega.hat * orbits$X + delta.hat)) +
beta.hat*sin(omega.hat*orbits$X + delta.hat)^2)

# Approx conf int for omega and delta
omega.hat + 1/sqrt(I.omega) * qnorm(0.975) * c(-1, 1)
delta.hat + 1/sqrt(I.delta) * qnorm(0.975) * c(-1, 1)

### (c)
### Bootstrap confidence intervals
B = 5000
boot.par = matrix(0, B, 2) # 2 columns for omega and delta
for (b in 1:B){
  # Generate a bootstrap sample (replace orbits$Y because of how I coded f3)
  orbits$Y = rnorm(n, alpha.hat + beta.hat * cos(omega.hat * orbits$X + delta.hat),
    sqrt(sig2.hat))
  # Compute mles for the bootstrap sample, with the actual mles as starting points
  # Not doing the iterative profile likelihood as before since it takes too long
  # and we are already near the correct values so there shouldn't be any issues
  temp = optim(c(alpha.hat, beta.hat, omega.hat, delta.hat, sig2.hat),
    function(x) -f3(x, log = TRUE), method = "L-BFGS-B",
    lower = c(min(orbits$Y), min(orbits$Y), 0, -pi, 0.001),
    upper = c(max(orbits$Y), max(orbits$Y), 3, pi, var(orbits$Y)))
  boot.par[b,] = temp$par[c(3,4)]
}

# Comparison
rbind(omega.hat + 1/sqrt(I.omega) * qnorm(0.975) * c(-1, 1),
  quantile(boot.par[,1], c(0.025, 0.975)))

rbind(delta.hat + 1/sqrt(I.delta) * qnorm(0.975) * c(-1, 1),
  quantile(boot.par[,2], c(0.025, 0.975)))

### Problem 2
### (b)
life = read.table("~/files/data/205b/lifetime.txt", header = TRUE)
n = nrow(life)
tobs = mean(life$x) - mean(life$y)
lhat = 0.5*(mean(life$x) + mean(life$y))
I.lambda = -(2*n/(lhat^2) - 2/(lhat^3) * (sum(life$x) + sum(life$y)))
pnorm(tobs / (1 / sqrt(I.lambda)),0,1,lower.tail = FALSE)

### (c)
### Permutation test
t.perm = double(10000)
w = as.double(unlist(life))
for (i in 1:length(t.perm)){
  s = sample(w)
  t.perm[i] = mean(head(s, 20)) - mean(tail(s, 20))
}
mean(t.perm >= tobs)

```