

Modeling Thefts in California Counties in 2010

Mickey Warner
AMS 207

Abstract

For the year 2010, counts of four types of theft (robbery, burglary, larceny, and motor vehicle) were recorded for 39 counties in California. The log-transformation is applied to the counts and the log-counts are modeled hierarchically as a multivariate normal of dimension four. We show that the multivariate normal provides a reasonable fit to the log-counts and report a summary for the predictive distributions for each theft by county.

1. Introduction

The data set we analyze in this paper contains the number of thefts (either robbery, burglary, larceny, or motor vehicle) for 39 California counties in 2010. The population for each county is also given, though we do not use this information in the analysis. We would expect that as population increases, the number of thefts also increase. By plotting population against each theft (not shown), we see that the relationship between population and theft count may be reasonably modeled under a multivariate regression framework. However, even with a log-transformation, the regression approach would have difficulty handling the heteroskedasticity that is present. We suspect there would also be issues in modeling those very low or very high population counties.

An alternative approach, the one we take in this paper, is to model the log theft counts hierarchically and to omit the population altogether. Figure 1 presents scatter plots of the four theft types for the original data (left) and the log-transformed data (right). The log transformation seems to suggest that a multivariate normal model is suitable. It is true that higher population counties observed greater numbers of thefts, but the hierarchical setting allows us the flexibility to dealing with any peculiarities. For instance, Sacramento county had an unusually low number of motor vehicle thefts considering it ranked near the top in all other thefts and is one of the most populated counties. The hierarchical model would account for this, whereas we could run the risk of overfitting in the regression case.

In section 2 we describe the details of the model and the model fitting procedure. A posterior analysis is given in section 3. We conclude with a discussion in section 4.

2. Methods

2.1 Hierarchical model

Denote \mathbf{z}_i as the vector of length $p = 4$ containing the log theft counts for county $i = 1, \dots, n = 39$. We assume a normal likelihood for each \mathbf{z}_i

$$\mathbf{z}_i | \boldsymbol{\mu}_i, \Sigma \sim N(\boldsymbol{\mu}_i, \Sigma) \quad (1)$$

and the following priors

$$\boldsymbol{\mu}_i | \boldsymbol{\mu}, V \sim N(\boldsymbol{\mu}, V) \quad (2)$$

$$\boldsymbol{\mu} \sim N(\mathbf{m}, C_0) \quad (3)$$

$$\Sigma \sim IW(S_0, r) \quad (4)$$

$$V \sim IW(D_0, k). \quad (5)$$

The likelihood assumes that the observation from each county comes from its own normal population. The prior on $\boldsymbol{\mu}_i$ (2) constrains the county means so we do not risk overfitting, which seems plausible given each observation has its own mean. We assume independence among $\boldsymbol{\mu}$, Σ , and V and that $\boldsymbol{\mu}_i$ are all independent.

Other prior specifications may model the data better, but these choices (2)-(5) result in convenient full posterior conditionals which we discuss later. The constants \mathbf{m} , C_0 , S_0 , r , D_0 , and k are chosen to be rather non-informative. With only $n = 39$ data points, we can not be as non-informative as we would like. We chose $\mathbf{m} = \bar{\mathbf{z}}$, $C_0 = I_p$, $S_0 = 0.25I_p$, $r = 6$, $D_0 = 3I_p$, and $k = 5$, where I_p is the $p \times p$ identity matrix. This reflects our expectation that the variance associated with each observation, Σ , should be “smaller” than the variance for the county means, V . Though not ideal, these priors were chosen after trial and error: they resulted in decent predictions while remaining somewhat non-informative.

2.2 Parameter estimation

As mentioned earlier, the prior specification results in convenient full conditionals. These are given by

$$\boldsymbol{\mu}_i | \cdot \sim N((\Sigma^{-1} + V^{-1})^{-1}(\Sigma^{-1}\mathbf{z}_i + V^{-1}\boldsymbol{\mu}), (\Sigma^{-1} + V^{-1})^{-1}) \quad (6)$$

$$\boldsymbol{\mu} | \cdot \sim N((nV^{-1} + C_0^{-1})^{-1}(nV^{-1}\bar{\boldsymbol{\mu}} + C_0^{-1}\mathbf{m}), (nV^{-1} + C_0^{-1})^{-1}) \quad (7)$$

$$\Sigma | \cdot \sim IW(S_0 + \sum_{i=1}^n (\mathbf{z}_i - \boldsymbol{\mu}_i)(\mathbf{z}_i - \boldsymbol{\mu}_i)^\top, r + n) \quad (8)$$

$$V | \cdot \sim IW(D_0 + \sum_{i=1}^n (\boldsymbol{\mu}_i - \boldsymbol{\mu})(\boldsymbol{\mu}_i - \boldsymbol{\mu})^\top, k + n) \quad (9)$$

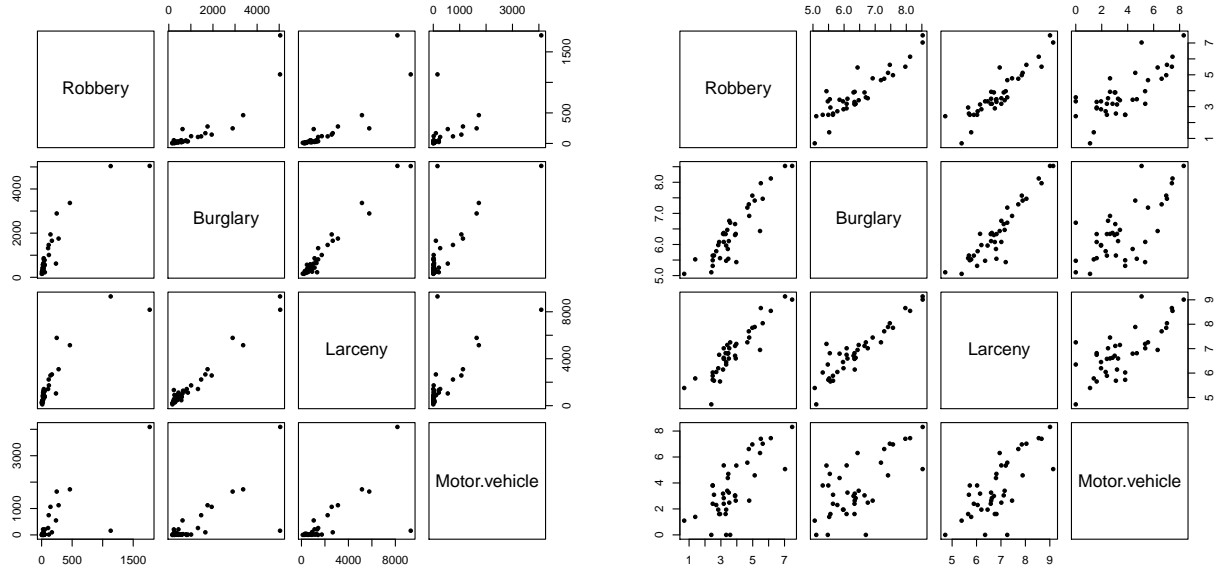


Figure 1: Scatter plots for the number of thefts by county. The left set of plots is the original data while the right set is of the log-transformed data.

where the dot (\cdot) represents the data and all other variables and $\bar{\mu} = 1/n \sum_{i=1}^n \mu_i$. We are thus able to explore the posterior using direct sampling. We iteratively draw samples from each full conditional based on the most recent samples of the other parameters.

3. Results

We obtain 10000 posterior draws after burning in 2000 draws. There did not appear to be any convergence or “stickiness” issues with the sampler as often happens when estimating several variances. Posterior means for μ , Σ , and V are as follows

$$\begin{aligned} \mu &= (3.761, 6.404, 6.863, 3.607)^\top \\ \Sigma &= \begin{pmatrix} 0.096 & 0.032 & 0.031 & -0.004 \\ 0.032 & 0.066 & 0.028 & -0.008 \\ 0.031 & 0.028 & 0.065 & -0.004 \\ -0.004 & -0.008 & -0.004 & 0.127 \end{pmatrix} \\ V &= \begin{pmatrix} 1.912 & 1.079 & 1.184 & 2.241 \\ 1.079 & 0.856 & 0.774 & 1.330 \\ 1.184 & 0.774 & 0.991 & 1.553 \\ 2.241 & 1.330 & 1.553 & 4.603 \end{pmatrix} \end{aligned}$$

The labels of each component in μ are Robbery, Burglary, Larceny, and Motor.Vehicle. The same order applies for the covariance matrices. Posterior means for each μ_i will not be presented, instead tables of predictions based on μ_i are given later. The overall mean μ corresponds very closely to the mean log thefts of the data (only off by a few thousandths in each component). As expected, the Σ is “smaller” than V as it has a smaller variance in each diagonal element.

Since we have a hierarchical model, there are various ways to obtain posterior predictions. For a particular county, we could use its county mean, μ_i , to make predictions by drawing \mathbf{z}_i^* from a $N(\mu_i, \Sigma)$. For a new county, we would need to draw μ_i^* from a $N(\mu, V)$ and then draw \mathbf{z}_i^* from $N(\mu_i^*, \Sigma)$. The predictions we will look at are based on the posterior samples for μ_i , for $i = 1, \dots, 39$.

Predictive distributions are obtained for each of the 39 counties and for each type of theft. Figure 2 shows our predictions (on the log-scale) compared to the observations. Equal-tailed 95% posterior predictive probability intervals are represented by the vertical lines and black dots represent the mean predictions. Each black dot and line pair is for a particular county. The plot suggests that our model provides a reasonable fit to the data: the mean predictions are all very close to the line $y = x$ and every interval crosses the line. There is a slight tendency to overestimate lower theft counts and underestimate higher theft counts, but not so much to be concerned with.

After drawing a \mathbf{z}_i^* from $N(\mu_i, \Sigma)$ using the posterior samples for μ_i and Σ we exponentiate the draw to get back to the original scale $\mathbf{y}_i^* = \exp(\mathbf{z}_i^*)$. We can use these back-transformed values to make inferences based on the theft counts for each county. Tables 1 through 4 contain summaries of all of these posterior predictive distributions. As expected, our model can predict a county’s theft count fairly well. However, this has its drawbacks which we will discuss in the next section.

Of particular interest is the probability of observing a total number of at least 3000 thefts in 2010 in Santa Cruz county. Santa Cruz county corresponds to observation $i = 31$. Using the back-transformed posterior predictive draws \mathbf{y}_{31}^* , of which we have 10000, we take each vector

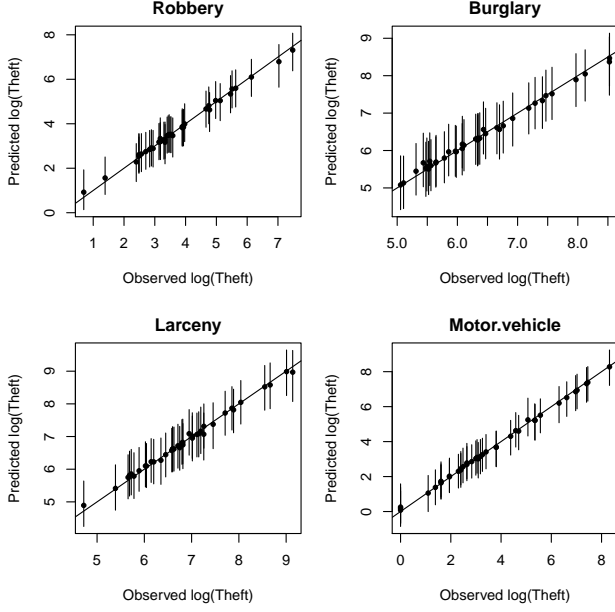


Figure 2: Plots of the posterior predictions (on the log-scale) against the observed log counts. The vertical lines represent equal-tailed 95% probability intervals and the solid dots are the mean predictions. The diagonal line is the line $y = x$.

and add up its components $\sum_{j=1}^4 y_{31,j}^*$. The probability of observing more than 3000 total thefts in Santa Cruz is computed by counting how many of these vectors had a sum greater than 3000 and then dividing that number by 10000. Our model estimates this probability at 0.0605.

4. Discussion

In this paper we fit a hierarchical model to multivariate normal data. The interest was mainly in obtaining posterior predictions for four theft types (robbery, burglary, larceny, and motor vehicle) in each of 39 counties in California. Despite Figure 2 indicating a good fit to the data, our model is rather limited. The model is really only useful in the year 2010. We would not expect to predict well for future years or for counties not included in the data set. Our predictions were good because we gave each county its own mean which is likely to change from year to year and certainly changes from county to county.

Note the predictive intervals from Tables 1-4. Every interval contains the observed value. This is more a sign of overfitting than it is of good model choice. More realistic predictions could be made by first drawing the mean μ_i and then drawing an observation \mathbf{z}_i as discussed in section 3. Doing so would increase the predictive variance, but the result is a more realistic prediction, especially for unobserved counties and years.

If future prediction was our primary concern, we would lean towards a regression model using population as a

Table 1: Summary of the predictive distributions for the number of robberies by county.

County	Obs	Mean	2.5%	50%	97.5%
Alameda	235	227	85	210	483
Butte	24	27	11	25	54
Calaveras	12	14	6	13	31
Contra Costa	119	111	39	105	220
El Dorado	34	37	14	34	76
Fresno	116	135	55	123	290
Humboldt	33	36	14	33	75
Imperial	13	16	6	14	35
Kern	463	483	186	453	971
Kings	12	15	6	13	36
Lake	17	19	8	17	40
Los Angeles	1770	1645	585	1565	3221
Madera	24	27	11	25	56
Marin County	28	26	9	24	53
Mendocino	19	20	8	18	41
Merced	49	52	21	49	105
Monterey	51	51	20	47	103
Napa	2	3	1	2	7
Nevada	4	5	2	5	12
Orange	32	38	15	33	83
Placer	30	35	14	32	74
Riverside	247	286	115	262	588
Sacramento	1131	988	293	951	1948
San Bernardino	145	169	67	153	360
San Diego	278	292	114	271	591
San Joaquin	168	168	64	157	337
San Luis Obispo	18	20	8	19	42
San Mateo	53	60	24	53	133
Santa Barbara	27	29	11	27	59
Santa Clara	24	31	13	27	74
Santa Cruz	36	35	12	33	72
Shasta	23	26	11	24	56
Sonoma	49	50	20	47	102
Stanislaus	106	115	47	106	237
Sutter	12	15	6	13	35
Tehama	11	11	4	10	23
Tuolumne	15	17	7	15	36
Ventura	31	36	15	32	80
Yuba	28	29	11	27	59

Table 2: Summary of the predictive distributions for the number of burglaries by county.

County	Obs	Mean	2.5%	50%	97.5%
Alameda	622	755	363	697	1498
Butte	564	568	263	542	1029
Calaveras	282	309	149	289	586
Contra Costa	1011	1014	452	966	1850
El Dorado	864	829	374	789	1505
Fresno	1466	1515	728	1430	2817
Humboldt	450	489	235	462	912
Imperial	283	316	154	292	620
Kern	3368	3285	1514	3139	5917
Kings	243	283	136	261	573
Lake	396	411	195	388	780
Los Angeles	5046	5056	2225	4822	9165
Madera	585	591	281	562	1078
Marin County	239	266	120	251	508
Mendocino	260	287	136	267	553
Merced	778	787	369	748	1436
Monterey	567	594	278	565	1082
Napa	157	171	83	157	343
Nevada	250	263	123	245	505
Orange	255	325	160	296	660
Placer	645	665	315	632	1232
Riverside	2892	2850	1316	2706	5219
Sacramento	5038	4624	1802	4472	8331
San Bernardino	1942	1959	919	1846	3596
San Diego	1755	1842	888	1744	3356
San Joaquin	1656	1618	729	1553	2913
San Luis Obispo	437	454	208	428	845
San Mateo	229	314	152	284	640
Santa Barbara	564	568	259	539	1039
Santa Clara	438	512	248	471	1034
Santa Cruz	812	756	306	725	1391
Shasta	568	575	272	542	1074
Sonoma	547	583	280	551	1088
Stanislaus	1320	1327	631	1254	2470
Sutter	203	248	122	227	504
Tehama	165	181	86	168	352
Tuolumne	325	350	169	329	666
Ventura	349	415	207	383	812
Yuba	390	419	200	395	785

Table 3: Summary of the predictive distributions for the number of larceny cases by county.

County	Obs	Mean	2.5%	50%	97.5%
Alameda	1041	1278	622	1184	2478
Butte	721	755	357	720	1371
Calaveras	361	407	198	378	778
Contra Costa	1729	1686	754	1612	3072
El Dorado	1114	1101	508	1053	1975
Fresno	2230	2383	1165	2243	4425
Humboldt	732	789	385	742	1465
Imperial	295	357	177	330	710
Kern	5152	5282	2513	5043	9580
Kings	306	375	180	343	743
Lake	492	532	254	502	987
Los Angeles	8174	8467	3731	8074	15764
Madera	755	792	386	751	1450
Marin County	573	562	249	537	1030
Mendocino	286	334	163	312	629
Merced	1227	1234	577	1173	2249
Monterey	735	793	381	753	1445
Napa	219	239	116	222	472
Nevada	324	349	167	326	673
Orange	910	978	474	904	1896
Placer	1275	1259	588	1187	2311
Riverside	5776	5617	2627	5365	10236
Sacramento	9309	8430	3360	8194	15004
San Bernardino	2574	2773	1331	2615	5246
San Diego	3106	3294	1563	3118	6090
San Joaquin	2664	2624	1232	2516	4671
San Luis Obispo	853	828	384	787	1516
San Mateo	1332	1388	664	1290	2637
Santa Barbara	912	897	421	858	1636
Santa Clara	1113	1185	575	1103	2278
Santa Cruz	1419	1258	526	1211	2294
Shasta	466	541	263	506	1044
Sonoma	821	874	423	829	1595
Stanislaus	1421	1585	759	1492	2949
Sutter	413	474	229	435	945
Tehama	112	143	70	132	289
Tuolumne	422	472	230	440	901
Ventura	894	960	463	900	1809
Yuba	634	665	318	630	1220

covariate. With some extra work, we could also add a spatial component to the model. Such a framework is likely to improve the predictive power over that of the hierarchical model. The hierarchical model does excel when we know a county's mean theft count, but this is a narrow situation.

Table 4: Summary of the predictive distributions for the number of motor vehicle thefts by county.

County	Obs	Mean	2.5%	50%	97.5%
Alameda	549	557	162	505	1283
Butte	11	13	5	11	31
Calaveras	11	12	4	11	28
Contra Costa	14	19	6	16	50
El Dorado	12	15	5	13	36
Fresno	746	750	229	693	1623
Humboldt	26	29	10	26	67
Imperial	22	23	7	21	53
Kern	1726	1825	576	1659	4040
Kings	45	45	13	41	99
Lake	7	8	3	7	19
Los Angeles	4090	4494	1412	4008	10277
Madera	17	20	7	17	47
Marin County	1	1	0	1	4
Mendocino	5	6	2	5	14
Merced	21	25	9	22	62
Monterey	14	17	6	15	42
Napa	3	3	1	3	8
Nevada	4	4	2	4	11
Orange	110	112	34	102	252
Placer	30	34	11	30	80
Riverside	1642	1683	511	1541	3793
Sacramento	159	225	75	179	647
San Bernardino	1064	1054	318	974	2263
San Diego	1124	1163	363	1065	2573
San Joaquin	98	116	40	101	286
San Luis Obispo	5	6	2	5	15
San Mateo	209	210	59	192	466
Santa Barbara	5	6	2	5	17
Santa Clara	211	208	58	190	461
Santa Cruz	1	2	1	1	5
Shasta	24	26	9	24	62
Sonoma	20	23	8	21	55
Stanislaus	261	277	93	251	620
Sutter	45	45	13	41	97
Tehama	1	1	0	1	3
Tuolumne	10	11	4	10	26
Ventura	80	83	25	75	186
Yuba	7	9	3	7	21