

Spectral analysis

Let $z_t = y_{t+1} - y_t$ for $t = 1, \dots, T - 1$ be the $T - 1$ first differences. The original data y_t is shown on the left side of Figure 1 and the first differences z_t are shown on the right side of Figure 1. We fit the model

$$z_t = a \cos(\omega t) + b \sin(\omega t) + \epsilon_t, \quad t = 1, \dots, T - 1$$

with $\epsilon_t \sim N(0, v)$. Under the prior $p(a, b, v, \omega) \propto v^{-1} p(\omega)$, we evaluate the marginal posterior for ω at the Fourier frequencies $\omega_k = 2\pi k/T$, for $1 \leq k < T/2$, as described in Prado and West (p. 87–88).

Figure 2 shows $p(\omega|\mathbf{z})$ for $\omega = \omega_k$ (left). This graph shows which frequencies are most prevalent in the data. Being on the log scale, it is clear that an angular frequency of about 1 is the dominating frequency. This corresponds to a wavelength (or period) $\lambda = 2\pi/\omega$ of about 6, or half a year.

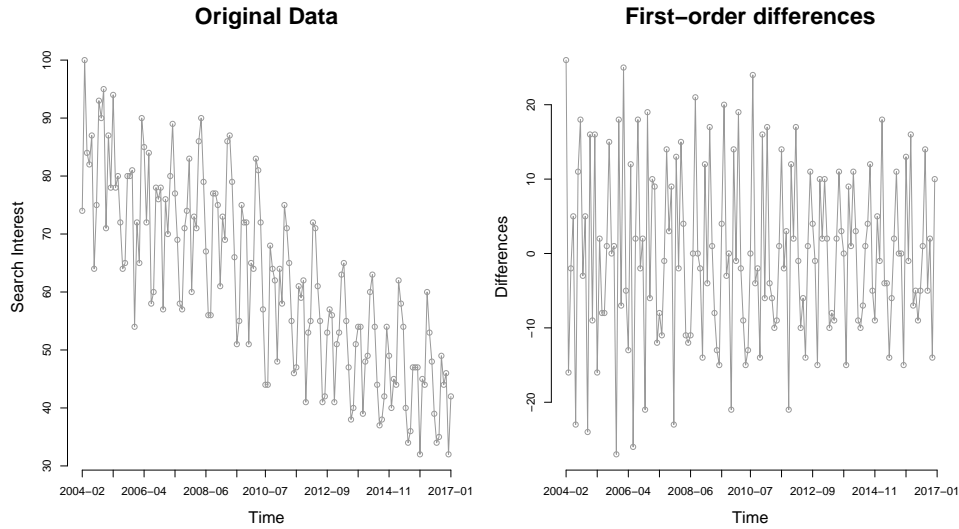


Figure 1: Left: plot of the data. Right: plot of the first differences of the data.




Figure 2: Left: log marginal posterior of the angular frequency $p(\omega|\mathbf{z})$. Right: log marginal posterior of the period/wavelength $p(\lambda = 2\pi/\omega|\mathbf{z})$. The data used is the first differences (right side of Figure 1).

Polynomial trend model

The DLM $\{1, 1, v, vW_t\}$, with v unknown, is fit to the original data $\mathbf{y} = (y_1, \dots, y_T)^\top$. The model is fit using the recursive formulas for normal DLMs. We integrate out v at each step so that $y_t|D_{t-1}$ and $\theta_t|D_t$ follow t - instead of normal distributions. The recursive formulas are given in West and Harrison (1997) and Prado and West (2010).

We use discount factors so that $W_t = (1 - \delta)/\delta C_{t-1}$, where C_{t-1} is the prior scale for $\theta_{t-1}|D_{t-1}$. δ is selected by maximizing the observed predictive density

$$p(\mathbf{y}|D_0) = \prod_{t=1}^T p(y_t|D_{t-1})$$

We evaluate $\log p(\mathbf{y}|D_0)$ on the grid $\delta = \{0.500, 0.501, \dots, 1.000\}$. The plot of δ vs. $\log p(\mathbf{y}|D_0)$ is given in Figure 3. The vertical line marks where the maximum occurs, at $\hat{\delta} = 0.867$, which is chosen to be our discount factor for the remainder of the analysis.




Figure 3: The log observed predictive density evaluated across several $\delta \in [0, 1]$. The dashed line indicates where the maximum occurs, at $\hat{\delta} = 0.867$.

We use the priors $m_0 = 60$, $C_0 = 200$, $n_0 = 1$, and $d_0 = 100$. The values on m_0 and C_0 correspond roughly to the sample mean and variance of \mathbf{y} , while the values for n_0 and d_0 are intended to be non-informative. Given $\delta = \hat{\delta} = 0.867$, we summarize the filtering distribution $\theta_t|D_t$, the one-step ahead forecast distribution $y_t|D_{t-1}$, and the smoothing distribution $\theta_t|D_T$, for $t = 1, \dots, T$, in Figures 4 and 5.

All three distributions appear to fit the data appropriately. The filtering and the forecasts capture the general downward trend of the data as well as the semiannual wavelength. And, as expected with smoothing, we capture the trend only. The variance for the one-step ahead forecast does seem a bit too high. The bands shown in the graph are 95% credible intervals, yet every observation falls within the bands.

Despite the positive appearance of these three distributions, when looking at the forecast distributions for θ_t and y_t (conditioned on D_T), the downward trend and the seasonality is no longer present (Figure 6). The trend might be captured by using time as a covariate in F_t , but we run the risk of specifying a feature in the model that may not exist in the data. We could incorporate seasonality into the model by specifying G_t to have a Fourier representation (P&W p.121), using the results from the earlier spectral analysis.




Figure 4: Left: mean of $\theta_t | D_t$ (thick green) with 95% credible intervals (thin green). Right: mean of the one-step ahead forecast function $y_t | D_{t-1}$ (thick blue) with 95% credible intervals (thin blue). Both are evaluated at $t = 1, \dots, T$.




Figure 5: Mean of smoothing distribution $\theta_t | D_T$ (thick red) with 95% credible intervals (thin red) at $t = T - 1, \dots, 1$.




Figure 6: Same as in Figure 4 except the darker shades indicate the forecast distributions $\theta_{T+h}|D_T$ (left) and $y_{T+h}|D_T$ (right) for $h = 1, \dots, 12$.